

SIMPLE ARTIFICIAL NEURAL NETWORK MODEL FOR CLASSIFICATION

```
In [1]: #importing relevant packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

In [2]: #reading the data
marks=pd.read_csv('Desktop/DATA/students_marks.csv')
```

```
In [3]: #head of data
marks.head()
```

	number_courses	time_study	Marks	Targets
0	3	4.508	19.202	0
1	4	0.096	7.734	0
2	4	3.133	13.811	0
3	6	7.909	53.018	1
4	8	7.811	55.299	1

```
In [4]: #dtypes
marks.dtypes
```

number_courses	int64
time_study	float64
Marks	float64
Targets	int64
dtype:	object

```
In [5]: #len of data
len(marks)
```

100

```
In [6]: #checking nulll values
marks.isnull().sum()
```

number_courses	0
time_study	0
Marks	0
Targets	0
dtype:	int64

```
In [7]: #train test split
```

```
In [8]: x=marks.drop('Targets',axis=1).values
y=marks['Targets'].values
```

```
In [9]: from sklearn.model_selection import train_test_split
```

```
In [10]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=101)
```

```
In [11]: y_test
```

```
Out[11]: array([1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1,
        1, 1, 0, 0, 1, 1, 1, 0, 1], dtype=int64)
```

```
In [12]: #data scalling
```

```
In [13]: from sklearn.preprocessing import MinMaxScaler
```

```
In [14]: scaler=MinMaxScaler()
```

```
In [15]: scaler.fit(x_train)
```

```
Out[15]: ▼ MinMaxScaler
MinMaxScaler()
```

```
In [16]: x_test=scaler.transform(x_test)
```

```
In [17]: x_train=scaler.transform(x_train)
```

```
In [18]: # MODEL TRAINING AND TESTING
```

```
In [19]: from tensorflow.keras.models import Sequential
```

```
In [20]: from tensorflow.keras.layers import Dense,Dropout
```

```
In [21]: x_train.shape
```

```
Out[21]: (70, 3)
```

```
In [22]: x_test.shape
```

```
Out[22]: (30, 3)
```

```
In [23]: y_test.shape
```

```
Out[23]: (30,)
```

```
In [24]: #model
```

```
In [25]: model=Sequential()
model.add(Dense(1000,activation ='relu'))
model.add(Dense(500,activation ='relu'))
model.add(Dense(500,activation ='relu'))
model.add(Dense(1,activation ='sigmoid'))
```

```
In [26]: model.compile(loss='binary_crossentropy',optimizer='adam',metrics='accuracy')
```

```
In [27]: from tensorflow.keras.callbacks import EarlyStopping
```

```
In [28]: early_stop=EarlyStopping(monitor='val_loss',mode='min',verbose=1,patience=25)
```


```
In [29]: #model fitting
```

```
In [30]: model.fit(x=x_train,y=y_train,epochs=20,batch_size=20,validation_data=(x_test,y_test),callbacks=[early_stop])
```

```
Epoch 1/20
4/4 [=====] - 1s 78ms/step - loss: 0.6622 - accuracy: 0.5429 - val_loss: 0.5604 - val_accuracy: 0.5667
Epoch 2/20
4/4 [=====] - 0s 27ms/step - loss: 0.5773 - accuracy: 0.5429 - val_loss: 0.4872 - val_accuracy: 0.6333
Epoch 3/20
4/4 [=====] - 0s 21ms/step - loss: 0.5194 - accuracy: 0.6714 - val_loss: 0.3952 - val_accuracy: 0.8667
Epoch 4/20
4/4 [=====] - 0s 23ms/step - loss: 0.4544 - accuracy: 0.8000 - val_loss: 0.3176 - val_accuracy: 0.9000
Epoch 5/20
4/4 [=====] - 0s 26ms/step - loss: 0.3899 - accuracy: 0.8286 - val_loss: 0.2553 - val_accuracy: 0.9333
Epoch 6/20
4/4 [=====] - 0s 21ms/step - loss: 0.3255 - accuracy: 0.8286 - val_loss: 0.1915 - val_accuracy: 1.0000
Epoch 7/20
4/4 [=====] - 0s 26ms/step - loss: 0.2817 - accuracy: 0.8714 - val_loss: 0.1424 - val_accuracy: 1.0000
Epoch 8/20
4/4 [=====] - 0s 26ms/step - loss: 0.2121 - accuracy: 0.9143 - val_loss: 0.1324 - val_accuracy: 0.9333
Epoch 9/20
4/4 [=====] - 0s 21ms/step - loss: 0.2096 - accuracy: 0.9143 - val_loss: 0.0799 - val_accuracy: 1.0000
Epoch 10/20
4/4 [=====] - 0s 26ms/step - loss: 0.1911 - accuracy: 0.8857 - val_loss: 0.0559 - val_accuracy: 1.0000
Epoch 11/20
4/4 [=====] - 0s 26ms/step - loss: 0.1357 - accuracy: 0.9429 - val_loss: 0.0972 - val_accuracy: 0.9667
Epoch 12/20
4/4 [=====] - 0s 25ms/step - loss: 0.1722 - accuracy: 0.9286 - val_loss: 0.0558 - val_accuracy: 1.0000
Epoch 13/20
4/4 [=====] - 0s 21ms/step - loss: 0.1460 - accuracy: 0.9143 - val_loss: 0.0622 - val_accuracy: 0.9667
Epoch 14/20
4/4 [=====] - 0s 26ms/step - loss: 0.1333 - accuracy: 0.9143 - val_loss: 0.0465 - val_accuracy: 1.0000
Epoch 15/20
4/4 [=====] - 0s 26ms/step - loss: 0.1399 - accuracy: 0.9429 - val_loss: 0.0448 - val_accuracy: 1.0000
Epoch 16/20
4/4 [=====] - 0s 26ms/step - loss: 0.1501 - accuracy: 0.9429 - val_loss: 0.0426 - val_accuracy: 1.0000
Epoch 17/20
4/4 [=====] - 0s 26ms/step - loss: 0.1069 - accuracy: 0.9429 - val_loss: 0.0858 - val_accuracy: 0.9667
Epoch 18/20
4/4 [=====] - 0s 24ms/step - loss: 0.1542 - accuracy: 0.9286 - val_loss: 0.0335 - val_accuracy: 1.0000
Epoch 19/20
4/4 [=====] - 0s 26ms/step - loss: 0.1278 - accuracy: 0.9286 - val_loss: 0.0339 - val_accuracy: 1.0000
Epoch 20/20
4/4 [=====] - 0s 21ms/step - loss: 0.1224 - accuracy: 0.9429 - val_loss: 0.0416 - val_accuracy: 1.0000
Out[30]: <keras.src.callbacks.History at 0x278b06d2d0>
```

```
In [31]: losses=pd.DataFrame(model.history.history)
```

```
In [32]: losses.plot()
```



```
In [33]: pred=(model.predict(x_test)>0.5).astype('int32')
```

```
1/1 [=====] - 0s 79ms/step
```

```
In [34]: from sklearn.metrics import mean_squared_error
```

```
In [35]: mean_squared_error(y_test,model.predict(x_test))
```

```
1/1 [=====] - 0s 32ms/step
0.009201021513787471
```

```
In [36]: np.sqrt(mean_squared_error(y_test,model.predict(x_test)))
```

```
1/1 [=====] - 0s 34ms/step
0.09592195532716934
```

```
In [37]: # model evaluation
```

```
In [38]: from sklearn.metrics import r2_score
```

```
In [39]: print(r2_score(y_test,pred))
```

```
1.0
```

```
In [40]: #checking with random forest classifier
```

```
In [41]: from sklearn.metrics import accuracy_score
```

```
In [42]: from sklearn.ensemble import RandomForestClassifier
```

```
In [43]: rfc=RandomForestClassifier()
```

```
In [44]: rfc.fit(x_train,y_train)
```

```
Out[44]: ▼ RandomForestClassifier
RandomForestClassifier()
```

```
In [45]: p=rfc.predict(x_test)
```

```
In [46]: print(r2_score(y_test,p))
```

```
0.8642533936651583
```

CONCLUSION

Artificial neural network was best with accuracy of 100% whereas Random forest classifier had 86% accuracy.